



Across The Globe

Named Entity Recognition (NER) and Feature Engineering

Analyzing Popularity and Accuracy with Random Forest

Novel Kathor

Date of Submission (06/01/2025)

Table of Contents

No	Objectives	Page No
01	Overview and Goals	01
02	Data Collection: Utilize the provided dataset of news articles.	02-03
04	Text Preprocessing: Conduct necessary preprocessing steps on the text.	04-05
04	Named Entity Recognition (NER) and Generating Synthetic Engagement Metrics	06-08
05	Feature Engineering: Combine entity counts with additional features to create a comprehensive feature set and Balanced the dataset.	10-15
06	Predictive Modeling: Training and evaluate predictive model	16-20
07	Visualization	21-29
08	Methodology for data preprocessing and feature extraction	30-34

09	Details about the predictive modeling process and performance metrics.	35-38
10	Insights on how named entities impact article engagement and popularity.	39-41

OBJECTIVE

The Objectives of this task is to analyze a set of news articles through Named Entity Recognition (NER) and to engineer numerical features based on these entities for predictive modeling. You will extract named entities from the articles using an open-source LLM model, create insightful features, and build a predictive model to determine article popularity based on engagement metrics. This task assesses your skills in natural language processing, feature engineering, and predictive analytics.

Goals:

1. Train a predictive model using the engineered features to predict article popularity.
2. Create visualizations to show the relationship between named entities and article popularity
3. Use csv id, url, title, tweet_ids to first figure out the news popularity by NER additionally showcase correlation between popular news with fake/real news.
4. Include relevant plots, such as:
 - Bar charts for entity frequency.
 - Scatter plots illustrating correlations.
 - Heatmaps for the relationship between entity counts and engagement metrics.

Data Collection

The dataset used for this project was sourced from the [FakeNewsNet](#) repository. This dataset is a comprehensive and publicly available resource widely used for research in fake news detection.

Dataset Overview

The FakeNewsNet dataset consists of news articles and associated metadata categorized into two domains:

1. GossipCop: A dataset focused on entertainment and celebrity news.
2. PolitiFact: A dataset centered around political news.

Files in the Dataset

The dataset is divided into four files based on the news source and authenticity:

1. gossipcop_fake.csv: Contains fake news articles from GossipCop.
2. gossipcop_real.csv: Contains real news articles from GossipCop.
3. politifact_fake.csv: Contains fake news articles from PolitiFact.
4. politifact_real.csv: Contains real news articles from PolitiFact.

Features

Each file contains the following features (columns):

- id: A unique identifier for each news article.
- news_url: The URL of the news article.
- title: The headline or title of the news article.
- tweet_ids: A list of tweet IDs associated with the news article, representing social media interactions related to the article.

Dataset Statistics

File	News Type	Number of Records
gossipcop_fake.csv	Fake	5336
gossipcop_real.csv	Real	16819
politifact_fake.csv	Fake	474

politifact_real.csv	Real	798
---------------------	------	-----

Data Sources

- GossipCop: Provides articles related to celebrity news, with information sourced from fact-checking by GossipCop.
- PolitiFact: Includes articles on political news, verified by PolitiFact's fact-checking team.

Purpose of Data Collection

The data collected from these files serves as the foundation for:

- Training machine learning models to classify fake and real news.
- Analyzing correlations between news popularity and authenticity.
- Extracting meaningful insights to combat misinformation.

Data Loading

The dataset files were loaded into Python using the `pandas` library, allowing for efficient manipulation and preprocessing. Each file was explored for data quality and completeness before analysis.

Challenges in Data Collection

1. Diversity of Sources: GossipCop and PolitiFact focus on different news domains, which may require tailored preprocessing steps.
2. Sparse Data in `tweet_ids`: Not all articles have associated social media interactions, requiring careful handling during analysis.
3. Class Imbalance: The uneven distribution of fake and real news articles introduces a challenge that needs to be addressed during preprocessing.

This structured dataset provides a robust basis for building machine learning models and performing in-depth analyses related to fake news detection.

Text Preprocessing

To ensure the dataset is clean and ready for analysis, several preprocessing steps were applied to the text data. This section outlines the steps taken to prepare the dataset for machine learning.

2.1 Cleaning and Standardizing Text

The following operations were performed on the title column of the dataset using the specified libraries:

- Removing Unnecessary Elements:
 - Library: re (Regular Expressions)
 - HTML tags, special characters, and extra whitespace were removed to eliminate noise.
- Normalizing Text:
 - Library: None (Standard Python string methods)
 - All text was converted to lowercase to maintain uniformity.
- Tokenization:
 - Library: nltk (Natural Language Toolkit)
 - The text was split into individual words (tokens) for further processing.
- Stop Word Removal:
 - Library: nltk (Stopword List)
 - Common words that do not contribute significant meaning (e.g., "the," "and") were removed to focus on informative terms.

2.2 Handling Missing and Irregular Data

Additional steps were performed to ensure the quality of the dataset using the following libraries:

- Removing Rows with Missing or Empty Cells:
 - Library: pandas
 - Rows with missing or empty cells were dropped.
- Filtering Excessively Large Data Entries:
 - Library: pandas
 - Rows with excessively large cell data were excluded based on a specified threshold.
- Validating and Filtering tweet_ids:
 - Library: pandas
 - Entries with invalid or missing tweet_ids were removed.
- Renaming Columns:
 - Library: pandas
 - The id column was renamed to news_id for better clarity and consistency.

2.3 Output of Preprocessing

The dataset was preprocessed and combined into a unified dataset with the following columns:

- news_id: Unique identifier for the news.
- news_url: URL of the news article.
- title: Original title of the news article.
- tweet_ids: Associated tweet IDs for the news.
- label: Label indicating whether the news is real or fake.
- cleaned_title: Preprocessed version of the title.

news_id	news_url	title	tweet_ids	label	cleaned_title
---------	----------	-------	-----------	-------	---------------

2.4 Final Dataset Details

After performing text preprocessing and filtering:

- Original number of rows: 23,196
- Rows kept: 13,110
- Rows removed: 10,086

The final combined and processed dataset contains a total of 13,111 rows, which will be used for further analysis and model training. The processed dataset is saved as dataset/filtered_data.csv.

Named Entity Recognition (NER)

Named Entity Recognition (NER) was performed on the cleaned titles to extract and categorize named entities. This step involves leveraging an open-source LLM model to identify entities such as organizations, locations, and people, enabling feature engineering for predictive modeling.

3.1 NER Process and Libraries Used

- NER Model: SpaCy (en_core_web_sm) was used for named entity extraction.
- Libraries Used:
 - spacy: To perform NER and extract entities.
 - pandas: For data manipulation and feature engineering.
 - tqdm: To monitor the progress of NER processing.

3.2 Extracting Entities and Features

- Entities Extracted:
 - PERSON: Named entities categorized as people.
 - ORG: Named entities categorized as organizations.
 - GPE: Named entities categorized as geopolitical entities (locations).
- Feature Engineering:
 - Entity counts for each category (person_count, org_count, gpe_count) were computed.
 - Names of extracted entities were stored in columns (person_name, org_name, locations_name).
 - Additional features such as title length, tweet word count, and sentiment polarity were added.

3.3 Output

The NER process enriched the dataset with new features and entity details, which can be used for further analysis.

4. Generating Synthetic Engagement Metrics

To simulate social media engagement, synthetic metrics for likes, shares, and comments were generated based on the extracted features and sentiment polarity.

4.1 Metric Generation and Libraries Used

- Libraries Used:
 - numpy: For generating random noise and performing numerical calculations.
 - pandas: To manipulate and enhance the dataset with synthetic metrics.
- Process:
 - A base engagement score was calculated using sentiment polarity, entity counts, and additional noise.
 - Metrics were derived as follows:
 - Likes: Base engagement + noise.
 - Shares: ~60% of likes + noise.
 - Comments: ~40% of likes + noise.

4.2 Output

The dataset was updated with synthetic engagement metrics (likes, shares, comments), providing realistic values for use in further modeling or analysis.

4.3 Final Processed File

The final processed dataset, including NER features and synthetic metrics, was saved as:

- NER Features: dataset/cleaned_data_with_entity_names.csv
- With Synthetic Metrics: dataset/cleaned_data_with_metrics.csv

Final Dataset: After performing entity extraction, sentiment analysis, and generating synthetic engagement metrics, the resulting dataset contains the following features:

- news_id: Unique identifier for the news article.
- news_url: The URL of the article.
- title: The original title of the article.
- tweet_ids: Tweet IDs associated with the article, if available.
- label: The target label, indicating the category or sentiment of the article.
- cleaned_title: The cleaned and pre-processed title.

- entities: A dictionary containing the extracted named entities (PERSON, ORG, GPE) and their names.
- person_count: The number of PERSON entities in the article.
- org_count: The number of ORG entities in the article.
- gpe_count: The number of GPE entities in the article.
- person_name: Comma-separated list of person names identified in the article.
- org_name: Comma-separated list of organization names identified in the article.
- locations_name: Comma-separated list of location names identified in the article.
- title_length: The length of the cleaned title.
- tweet_count: The number of words in the cleaned title.
- sentiment: The sentiment polarity of the cleaned title.
- likes: Synthetic metric for likes.
- shares: Synthetic metric for shares.
- comments: Synthetic metric for comments.

Feature Engineering: Combine entity counts with additional features to create a comprehensive feature

■ Entity Extraction:

- The code uses SpaCy to perform Named Entity Recognition (NER) on the `cleaned_title` of each news article. It extracts named entities of types such as PERSON, ORG (organizations), and GPE (geopolitical entities, e.g., locations). These entities are stored in a dictionary, which is later expanded into individual counts and lists of entity names.
- Entities Dictionary: For each article, a dictionary of the form {'PERSON': count, 'ORG': count, 'GPE': count, 'person_name': list, 'org_name': list, 'locations_name': list} is generated.
- This dictionary is stored in the `entities` column of the dataset, which provides a breakdown of entities for further analysis.

■ Feature Engineering:

- Person, Organization, and Location Counts: Based on the NER results, counts of persons (`person_count`), organizations (`org_count`), and locations (`gpe_count`) are computed. These counts provide insight into the article's content and focus.
- Person, Organization, and Location Names: The names of persons, organizations, and locations identified in the article titles are stored as comma-separated lists in `person_name`, `org_name`, and `locations_name` columns, respectively. This helps in identifying key subjects in the article.

■ Additional Features:

- Title Length: The length of the cleaned title (in characters) is computed and stored in the `title_length` column.
- Tweet Count: The number of words in the cleaned title is counted and stored in the `tweet_count` column. This feature can give insight into how concise or detailed the article title is.

■ Sentiment Analysis:

- The code uses TextBlob to calculate the sentiment polarity of the cleaned title. Sentiment polarity ranges from -1 (negative sentiment) to 1 (positive sentiment). This is stored in the `sentiment` column and gives an indication of the emotional tone of the article.

person_count	org_count	gpe_count	title_length	tweet_count	sentiment
1	0	1	76	10	0.50
2	0	0	66	10	0.30
1	0	0	47	7	0.20
1	1	0	56	9	0.00

■ Synthetic Engagement Metrics:

- The code generates synthetic engagement metrics (likes, shares, comments) based on a formula involving the sentiment score and entity counts. These metrics simulate how articles might be engaged with in the real world. For example:
 - Likes: A base engagement score is calculated using sentiment, person count, organization count, and location count, and then noise is added to simulate real-world variability.
 - Shares and Comments: These are derived as a percentage of likes, with some noise added for realism.
- These synthetic metrics are stored in the likes, shares, and comments columns.

likes	shares	comments
64	33	22
55	20	12
55	25	11
35	13	7
59	22	14

Balanced the dataset

What is Unbalanced Data?

Unbalanced data refers to a situation in machine learning where the distribution of the target classes is not uniform. In other words, one class significantly outnumbers the other(s), leading to a skewed dataset. This imbalance can cause issues in model training, as machine learning algorithms may be biased towards the majority class and fail to properly learn patterns for the minority class. In classification tasks, such as binary classification (e.g., predicting real vs. fake news), the model might end up predicting the majority class most of the time, because that class has more data and therefore dominates the learning process.

In the case of binary classification, for instance:

- Majority class: The class that has more samples.
- Minority class: The class that has fewer samples.

Example of Unbalanced Data:

Let's consider a dataset that classifies news articles as either real (label 1) or fake (label 0). If the number of real news articles is significantly higher than fake news articles, we would have an unbalanced dataset.

Before applying any balancing techniques, the distribution of labels in your dataset might look like this:

- Real (label 1): 9206 articles
- Fake (label 0): 3904 articles

In this case, there are 9206 real news articles and only 3904 fake news articles, making the dataset highly unbalanced. The ratio of real to fake news is approximately 2.4:1 (real to fake). This imbalance can result in the model becoming biased towards predicting the majority class (real news) more frequently.

How Much Data Do We Have Before Balancing?

Before balancing, the dataset contains the following:


- Total number of samples: 9206 (real news) + 3904 (fake news) = 13110 samples
- Real news (label 1): 9206 articles

- Fake news (label 0): 3904 articles

The class distribution is:

- Real news: 9206 samples (70.2%)
- Fake news: 3904 samples (29.8%)

This shows that the dataset is not balanced, with the real news class being larger than the fake news class.



the goal is to balance the dataset by addressing the class imbalance between the real (label 1) and fake (label 0) news articles. Here's a detailed explanation of how the code works to balance the dataset:

1. Initial Dataset Inspection:

- The dataset is loaded using pandas with the `load_and_prepare_data` function, which reads a CSV file containing the news articles and displays the initial distribution of labels (1 for real news and 0 for fake news).

2. Entity Parsing:

- The `parse_entities` function is used to parse the `entities` column in the dataset, which contains a dictionary of extracted named entities (e.g., PERSON, ORG, GPE). This function ensures that if the `entities` column has any malformed data, it returns a default dictionary with zero counts for each entity type.

3. Data Augmentation (Generating Variations):

- To balance the dataset, the code generates new samples for the minority class (label 0, fake news). This is done using the `generate_variations` function, which creates multiple variations of the fake news article's title by:
 - Entity Replacement: Replaces named entities like persons, organizations, and locations with other similar entities (e.g., replacing a PERSON entity with "John Smith").
 - Synonym Replacement: Replaces adjectives, nouns, and verbs with similar words from word vectors.
 - Word Order Modification: Shuffles the order of words in the title if the title is long enough.

4. Entity Update and Feature Engineering:

- Once a new variation of the article title is generated, the `update_entities` function is called to extract named entities from the new title. The function counts the occurrences of PERSON, ORG, and GPE entities and returns a dictionary with the entity counts and the actual names of the entities.
- Additionally, the code computes and updates other features like:
 - `person_count`: The number of PERSON entities in the title.
 - `org_count`: The number of ORG entities in the title.
 - `gpe_count`: The number of GPE entities in the title.
 - `person_name`, `org_name`, `locations_name`: Comma-separated lists of identified entities.

5. Dataset Balancing:

- The code calculates how many new samples are needed to balance the dataset by subtracting the number of minority class samples (label 0) from the majority class samples (label 1).
- New samples are generated by taking variations of the titles from the minority class (label 0) and adding them to the dataset. Each minority class sample can generate up to two variations, and the loop ensures that enough new samples are created to match the number of majority class samples.

6. Shuffling and Saving the Balanced Dataset:

- Once the new samples are generated, they are concatenated with the original dataset (`df`) to create a combined dataset with balanced classes (label 1 and label 0 having equal counts).
- The combined dataset is then shuffled to ensure that the data is randomly distributed.
- Finally, the balanced dataset is saved to a new CSV file (`dataset/resampled_data.csv`), and the final distribution of labels is printed, which should now show an equal number of real and fake news articles.

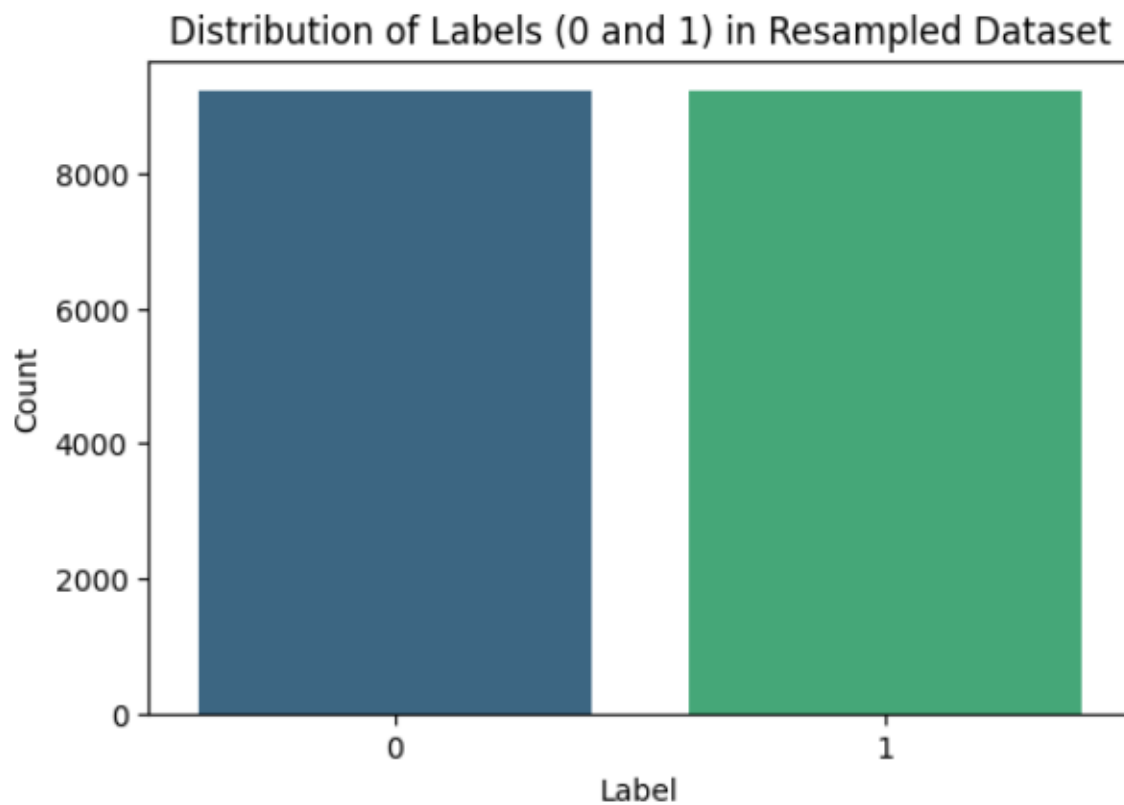
7. Final Label Distribution: After balancing the dataset, the final distribution of labels is:

```
yaml
Copy code
label
1  9206
0  9206
Name: count, dtype: int64
```

This ensures that there are 9206 samples of both real and fake news articles, achieving a balanced dataset.

Key Features of the Code:

- **Data Augmentation:** The code uses multiple techniques (entity replacement, synonym replacement, and word order modification) to generate new samples for the minority class.
- **Entity and Feature Engineering:** For each generated variation, the code extracts and updates features like entity counts and sentiment-related metrics.
- **Balanced Dataset:** The code generates enough new samples to ensure that both classes have equal representation, thus solving the problem of class imbalance.
- **Saving the Balanced Dataset:** The resulting balanced dataset is saved for further analysis or model training.



After balancing, the distribution is:

- Real news (label 1): 9206 samples
- Fake news (label 0): 9206 samples

Now, the dataset contains 18412 samples in total, with an equal number of real and fake news articles, which eliminates the bias in the model's learning process.

Balancing the dataset helps ensure that the model can learn the patterns and characteristics of both classes equally well. It prevents the model from being biased towards the majority class and improves its ability to generalize and correctly classify instances of the minority class (in this case, fake news). Balancing can be achieved through techniques like:

- Oversampling the minority class (e.g., by generating synthetic data or duplicating minority class samples).
- Undersampling the majority class (e.g., by removing some samples from the majority class).

Predictive Modeling: Training and evaluate predictive model

Predictive Modeling for Article Popularity Prediction

In this section, we focus on building a predictive model to predict article popularity using various engineered features from the dataset. Our goal is to develop a model that can predict whether an article will be considered popular or not based on its content and metadata.

1. Data Preparation

We begin by loading the dataset, which contains multiple columns representing different features of the articles, including the title, number of likes, shares, comments, and other metadata. After ensuring that any missing values in the columns for `person_name`, `org_name`, and `locations_name` are filled with a placeholder, we perform feature engineering.

2. Feature Engineering

We extract several features from the text data, such as:

- Sentiment Score: The sentiment of the article's title, derived using the TextBlob library.
- Title Length: The length of the article's title in characters.
- Word Count: The number of words in the article's title.

Additionally, we combine the `person_name`, `org_name`, and `locations_name` columns into a new feature called `combined_entities`, which serves as an additional textual feature.

3. Text Vectorization

We use the TF-IDF Vectorizer to transform the textual data (article title and combined entities) into numerical vectors that can be fed into the machine learning model. We limit the number of features to 5000, as a higher number of features may lead to overfitting.

4. Data Splitting

The dataset is split into training and testing sets, with 75% of the data used for training and 25% for testing. This ensures that we have a separate dataset for evaluation, allowing us to assess the model's generalization ability.

5. Model Selection

We use the Random Forest Classifier for this task due to its effectiveness in handling complex, high-dimensional data and its ability to handle both numerical and categorical features. The Random Forest algorithm builds multiple decision trees and combines their predictions, which generally leads to a more robust model.

6. Model Training and Hyperparameter Tuning

To optimize the model's performance, we perform grid search for hyperparameter tuning. We explore various combinations of hyperparameters, such as:

- `n_estimators` (number of trees in the forest)
- `max_depth` (maximum depth of each tree)
- `min_samples_split` (minimum number of samples required to split an internal node)
- `min_samples_leaf` (minimum number of samples required to be at a leaf node)

By iterating over all possible combinations of hyperparameters, we identify the best-performing set of parameters.

The code trains a model using the following features:

1. Text Features:

- `title`: The title of the article.
- `combined_entities`: A combination of `person_name`, `org_name`, and `locations_name`. This feature represents named entities such as persons, organizations, and locations within the text.

These text features are processed using a TF-IDF Vectorizer (`TfidfVectorizer`) to extract numeric features based on word frequencies, with a maximum of 5000 features.

2. Numerical Features:

- `sentiment_score`: The sentiment polarity of the article's title, calculated using `TextBlob`. This score reflects the sentiment of the text, with values ranging from -1 (negative) to 1 (positive).
- `title_length`: The length of the article's title in characters.
- `word_count`: The word count of the article's title.

These numerical features capture basic properties of the article's title, such as its sentiment, length, and word complexity.

3. Additional Data Features:

- likes: The number of likes the article received.
- shares: The number of shares the article received.
- comments: The number of comments the article received.

These features give insights into the article's engagement metrics, which are often associated with its popularity.

Data Preprocessing:

- Missing Values: Missing values in person_name, org_name, and locations_name are filled with the placeholder 'empty'.
- Text Feature Vectorization: The text data (title + combined_entities) is vectorized into numerical features using TF-IDF, which captures the importance of words in the context of the entire dataset.

Target Variable (y):

- label: This is the target variable representing whether the article is considered "popular" or not, based on its engagement metrics (likely derived from likes, shares, and comments).

Data Resampling:

- SMOTE (Synthetic Minority Over-sampling Technique) is used to handle class imbalance in the dataset by generating synthetic samples for the minority class during training.

These features are then used to train a Random Forest Classifier, with model evaluation done using accuracy, classification report, confusion matrix, and ROC-AUC score.

Model Evaluation Metrics¹

This section provides a comprehensive analysis of the performance of the classification model, which is used to predict news popularity and its classification as fake or real. Below are the evaluation metrics including the **Classification Report**, **Confusion Matrix**, and **ROC-AUC Score**.

1. Classification Report

The classification report provides several important metrics that evaluate the model's performance for both classes (0: Fake, 1: Real):

- **Precision:**

- For class 0 (Fake), the precision is **0.92**, meaning 92% of the predictions for fake news were correct.
- For class 1 (Real), the precision is **0.80**, indicating that 80% of the real news predictions were accurate.
- **Recall:**
 - For class 0 (Fake), recall is **0.76**, meaning 76% of the actual fake news instances were correctly identified.
 - For class 1 (Real), recall is **0.93**, meaning 93% of the actual real news instances were correctly identified.
- **F1-Score:**
 - For class 0 (Fake), the F1-score is **0.83**.
 - For class 1 (Real), the F1-score is **0.86**.
- **Accuracy:** The overall accuracy is **0.85**, meaning 85% of all predictions were correct.
- **Macro Average:**
 - Precision: **0.86**
 - Recall: **0.85**
 - F1-Score: **0.85**
- **Weighted Average:**
 - Precision: **0.86**
 - Recall: **0.85**
 - F1-Score: **0.85**

2. Confusion Matrix

The confusion matrix is as follows:

[[1721 550]

[154 2178]]

ROC-AUC Score: **0.9149**

Random Forest - **Train Accuracy: 0.8927**

Random Forest - **Test Accuracy: 0.8471**

Mean Absolute Error (MAE)

- The **Mean Absolute Error (MAE)** is **0.1529**, which is a measure of the average absolute difference between predicted and actual values. Since the MAE is relatively low, it indicates that the model's predictions are close to the true values on average.

Summary of Performance:

- **Accuracy:** 84.71% — Strong overall correctness of predictions.
- **F1-Score:** 86.09% — The model performs well in balancing both precision and recall.
- **MAE:** 15.29% — The model has a small average error, indicating it is fairly accurate in predicting the outcome.

These results suggest that the model is robust and performs well, making it a good candidate for tasks involving fake vs. real news classification, especially in terms of overall accuracy and balanced performance.


Visualization

Data Visualization of Most Frequent Entities

In this notebook, we visualize the top 10 most frequent entities from three different columns: `person_name`, `org_name`, and `locations_name`. The visualization helps us better understand the distribution of these entities in the dataset.

Key Steps:

1. **Data Filtering:** We remove rows where the `person_name`, `org_name`, or `locations_name` contain the value 'empty' to focus on meaningful entries.
2. **Entity Frequency Calculation:** We calculate the count of occurrences of the top 10 most frequent entries for each of the three columns.
3. **Data Visualization:** The counts of these top 10 entities are visualized using bar plots for better interpretation.



Popularity Analysis of Articles Based on Named Entities

In this notebook, we explore the relationship between the popularity of articles and three named entities: `person_name`, `org_name`, and `locations_name`. We define popularity as the sum of likes, shares, and comments and visualize the top 10 most popular articles based on each entity.


Steps:

1. **Load and Preprocess Data:**
 - The dataset is loaded and column names are cleaned by stripping any leading or trailing spaces.
 - A new metric, popularity, is calculated as the sum of likes, shares, and comments.
2. **Data Filtering:**

- Rows where the person_name, org_name, or locations_name columns have the value 'empty' are removed to ensure the integrity of the analysis.
3. Entity Popularity Calculation:
 - The function get_entity_popularity() calculates the average popularity for each entity by grouping the data based on the entity (person, organization, or location).
 4. Visualization:
 - Bar plots are generated to visualize the average popularity for the top 10 entities in each of the three categories: person_name, org_name, and locations_name.

Outcome:

The result is a set of visualizations showing the top 10 most popular entities across the three categories, helping us understand the impact of named entities on article popularity.



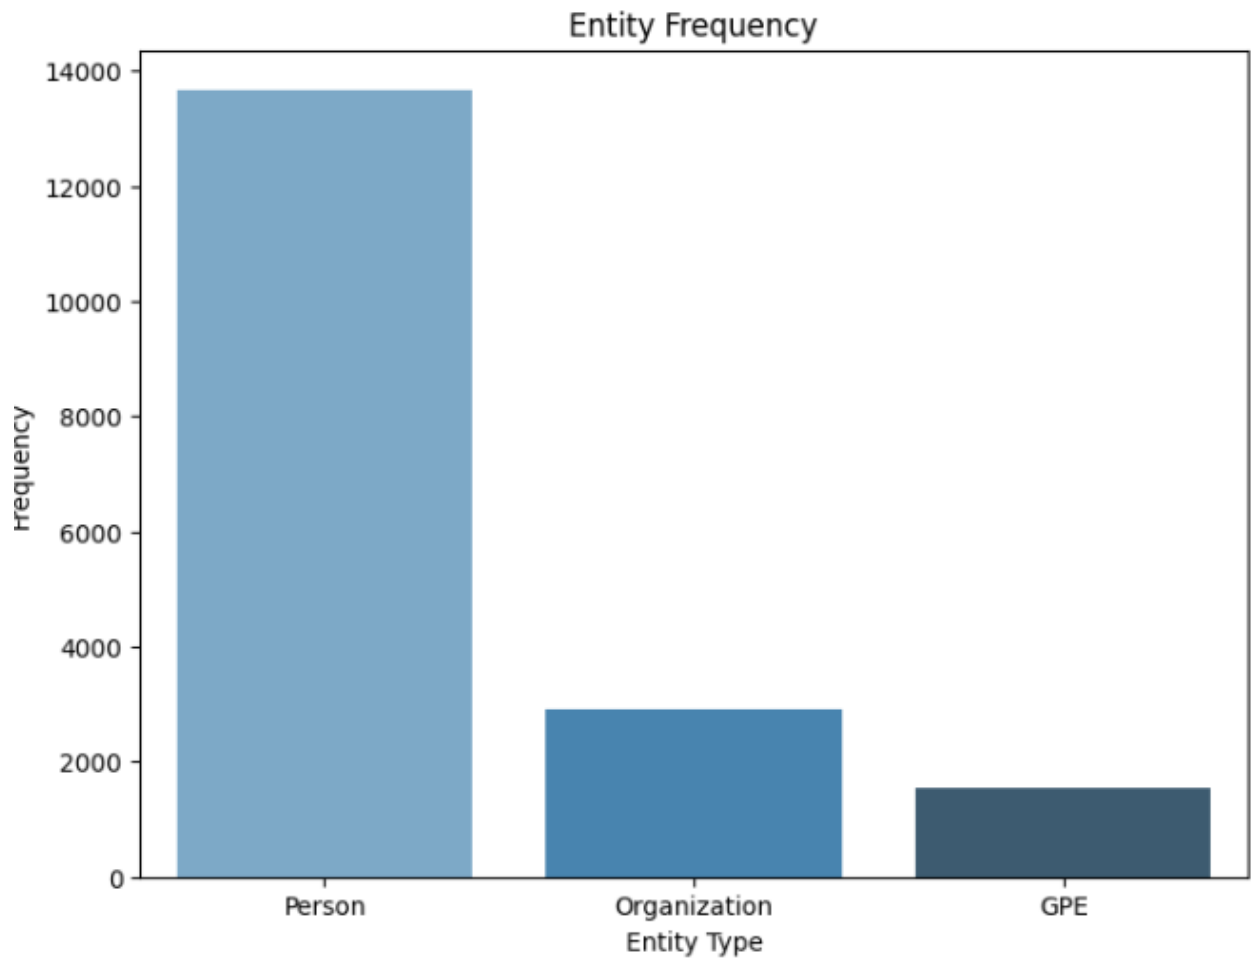
Entity Frequency Analysis

In this notebook, we visualize the frequency of different entity types (Person, Organization, and GPE) across the dataset. We aggregate the counts for each entity type and present the results in a bar plot. This helps in understanding the distribution of different named entities in the data.

Steps:

1. Data Loading:
 - The dataset is loaded from a CSV file (resampled_data.csv).
2. Data Aggregation:
 - The counts for each entity type—person_count, org_count, and gpe_count—are summed up to get the total frequency for each entity type.
3. Visualization:
 - A bar plot is generated using Matplotlib and Seaborn to visualize the frequency of each entity type. The x-axis represents the entity types (Person, Organization, and GPE), and the y-axis shows the total frequency for each type.

The plot provides a clear view of how many instances of each entity type are present in the dataset.



Scatter Plot Analysis: Entity Counts vs Engagement Metrics

In this section, we analyze the relationship between different entity counts (person, organization, and GPE) and engagement metrics (likes, shares, and comments). We visualize these relationships using scatter plots to identify potential trends or correlations.

Scatter Plot 1: Person Count vs Likes

This scatter plot shows the relationship between the count of person entities (person_count) and the number of likes (likes). Each point represents an observation from the dataset, with the x-axis showing the person count and the y-axis showing the corresponding number of likes.

Scatter Plot 2: Organization Count vs Shares

Here, we visualize the relationship between the count of organization entities (org_count) and the number of shares (shares). The x-axis shows the organization count, and the y-axis shows the number of shares associated with each observation.

Scatter Plot 3: GPE Count vs Comments

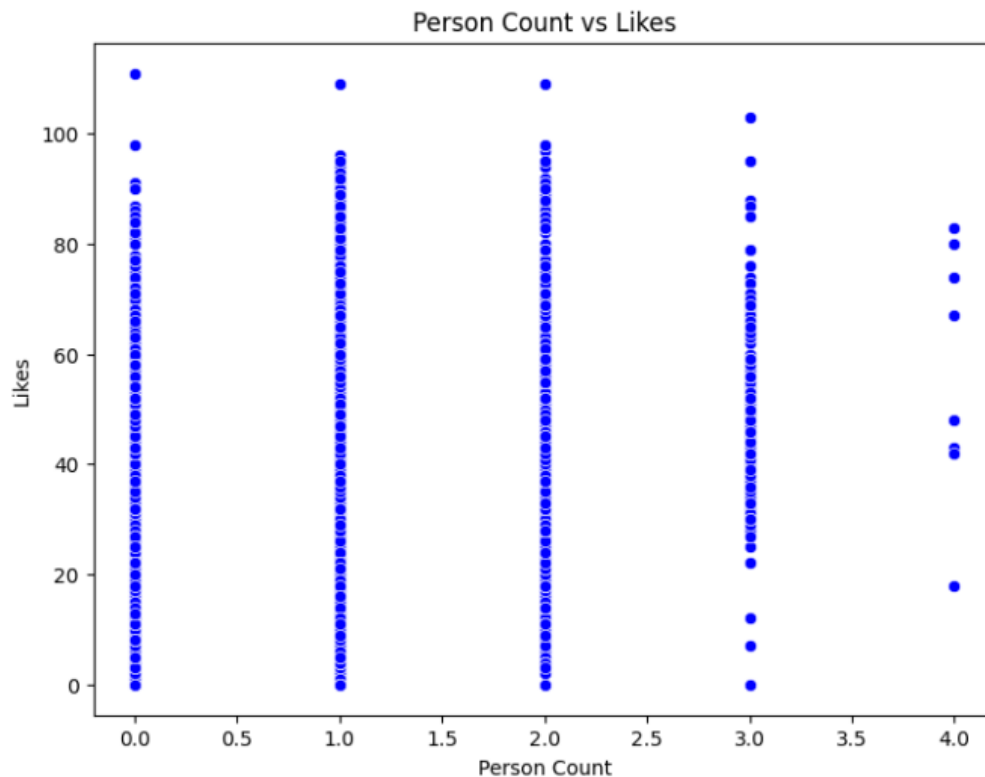
In this plot, we examine the relationship between the count of GPE entities (`gpe_count`) and the number of comments (`comments`). The x-axis represents the GPE count, while the y-axis represents the number of comments.

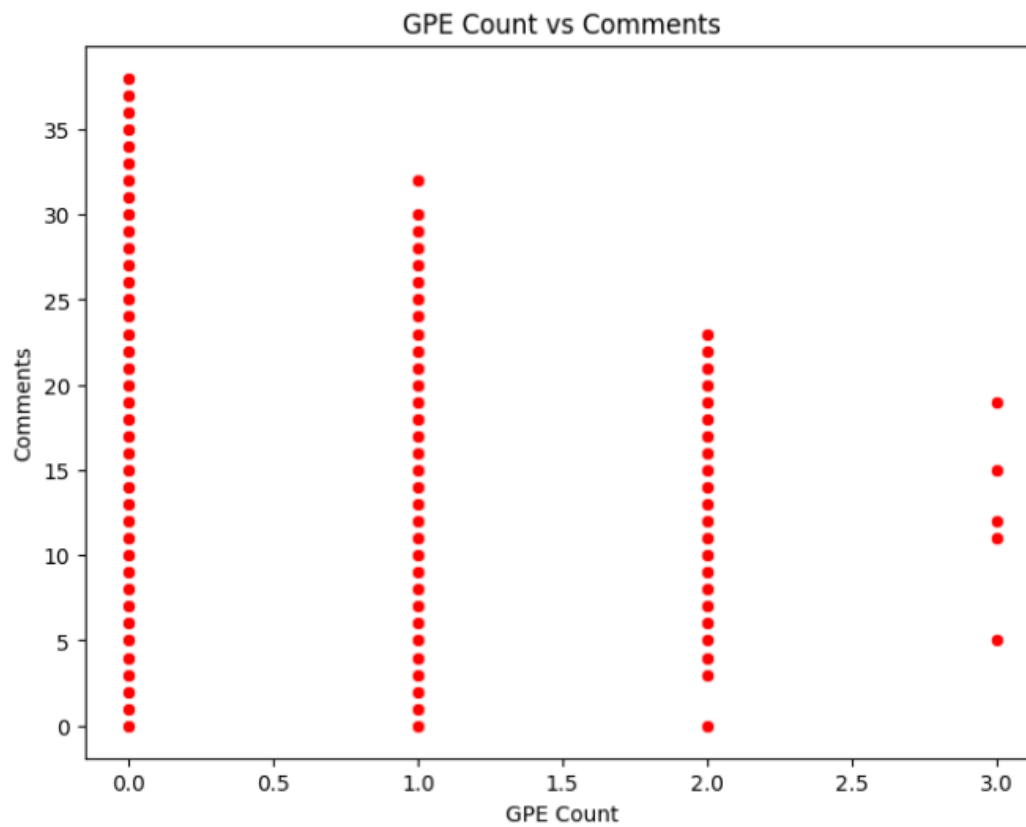
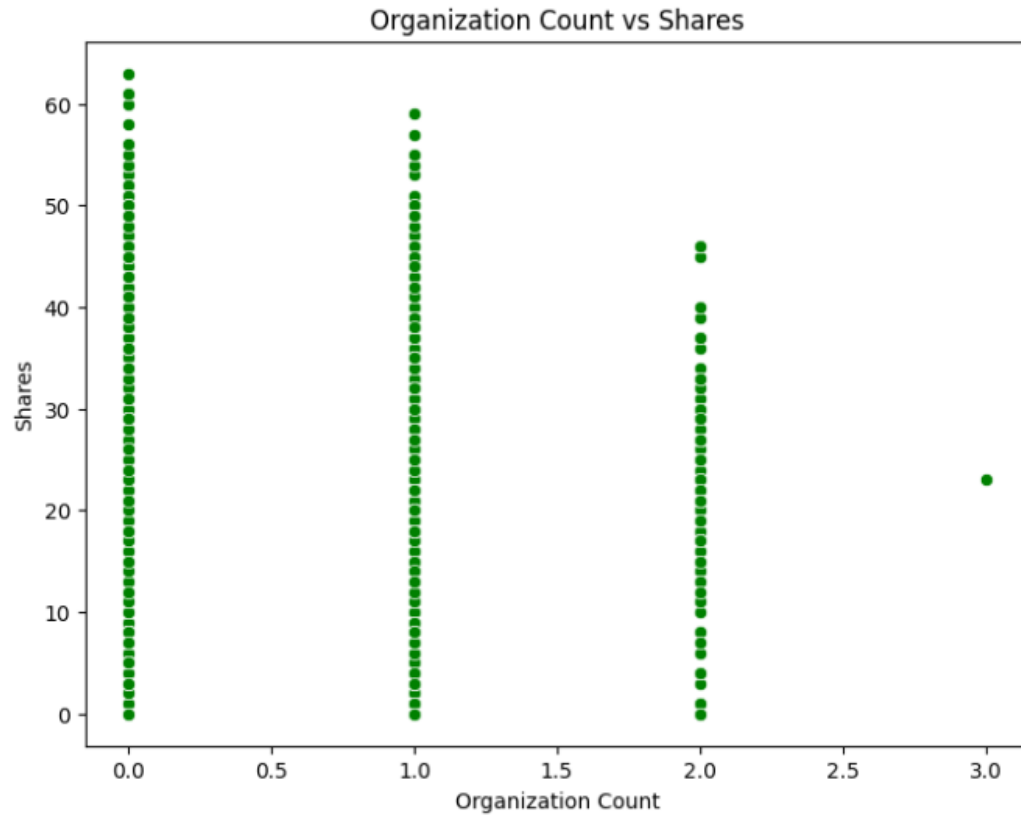
Steps:

1. Plotting Person Count vs Likes:
 - A scatter plot is generated with the `person_count` on the x-axis and `likes` on the y-axis. Each point represents a single row from the dataset.
2. Plotting Organization Count vs Shares:
 - A similar scatter plot is created for the `org_count` vs `shares`.
3. Plotting GPE Count vs Comments:
 - The final scatter plot visualizes the relationship between `gpe_count` and `comments`.

Insights:

These plots allow us to explore how the count of different entity types might be correlated with various forms of engagement (likes, shares, comments). Observing these relationships may help uncover potential trends or patterns within the data.





Correlation Heatmap: Entity Counts and Engagement Metrics

In this section, we analyze the relationships between different entity counts (person_count, org_count, gpe_count) and engagement metrics (likes, shares, comments) using a correlation heatmap.

Heatmap Overview:

A correlation heatmap is used to visualize the correlation coefficients between the selected variables. Correlation values range from -1 to 1, where:

- 1 indicates a perfect positive correlation.
- -1 indicates a perfect negative correlation.
- 0 indicates no correlation.

The heatmap allows us to quickly assess the strength of relationships between the counts of entities and the engagement metrics.

Code Breakdown:

1. Correlation Calculation:

- The correlation matrix is calculated using the `corr()` function, which computes pairwise correlation coefficients for the selected columns: `person_count`, `org_count`, `gpe_count`, `likes`, `shares`, and `comments`.


2. Heatmap Plot:

- A heatmap is created using Seaborn's `heatmap()` function. The correlation values are displayed within the cells using `annot=True`, and the color palette `coolwarm` is applied to visually indicate the strength of correlations.
- The `fmt='.2f'` argument formats the correlation values to two decimal places.

Insights:

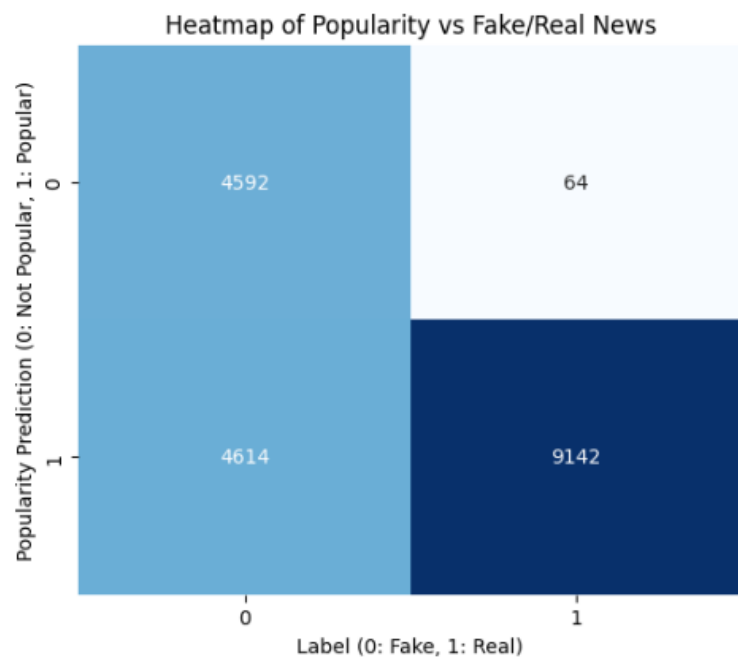
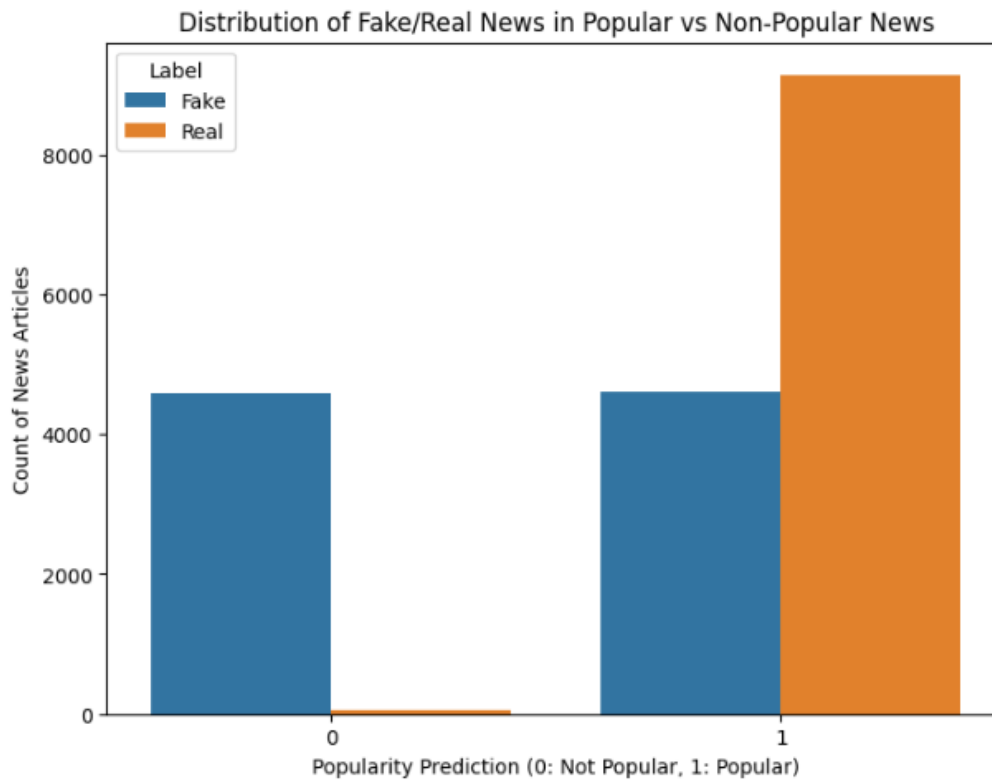
- Positive Correlations:
 - Strong positive correlations between entity counts (such as person, organization, and GPE) and engagement metrics might suggest that higher counts of entities lead to increased likes, shares, or comments.
- Negative Correlations:
 - Negative correlations would suggest that as one variable increases, the other tends to decrease.
- Weak or No Correlations:
 - A correlation close to 0 suggests there is little to no linear relationship between the variables.

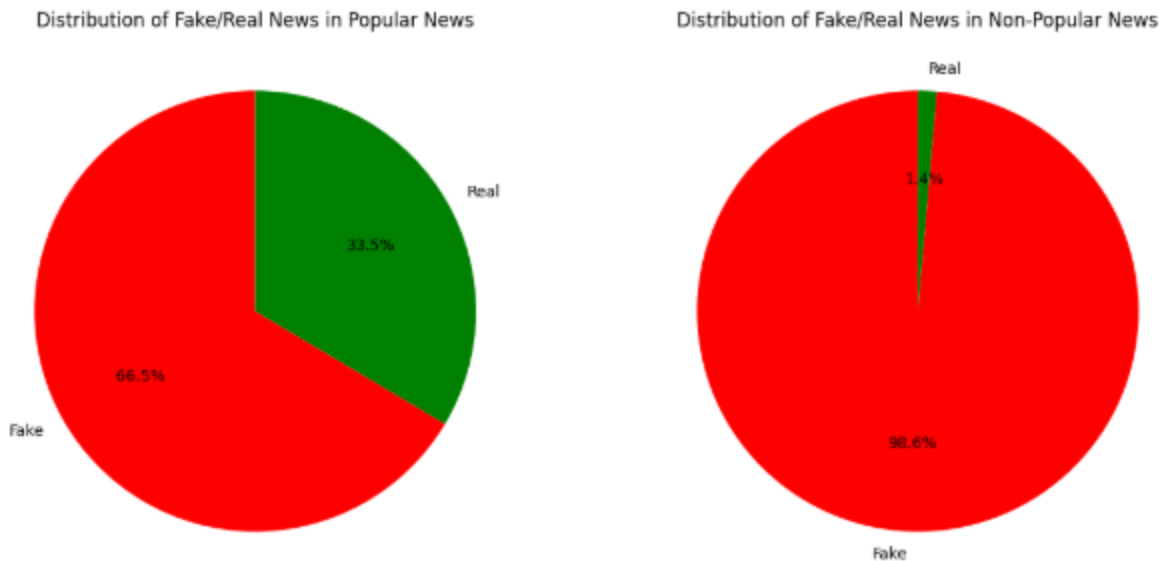
This heatmap serves as an effective tool to identify potential areas for further analysis and help understand how the entity counts are related to user interactions.



1. Bar Plot: Distribution of Fake/Real News in Popular vs Non-Popular News

The first visualization is a bar plot that shows the count of fake and real news articles, categorized by the popularity prediction. The popularity prediction is binary: 0 for non-popular and 1 for popular news.





Data Visualization for Fake/Real News Analysis

This section provides visual representations of the relationship between the predicted popularity of news articles and their classification as fake or real. The visualizations used include a count plot, heatmap, and pie charts. These plots help in understanding the distribution and correlation of fake and real news articles across popular and non-popular categories.

1. Bar Plot: Distribution of Fake/Real News in Popular vs Non-Popular News

The bar plot shows the count of fake and real news articles for each popularity category (Popular vs Non-Popular). The plot makes it easy to compare the distribution of fake and real news between these two categories.

2. Heatmap: Popularity vs Fake/Real News

This heatmap illustrates the relationship between the popularity of news articles and their classification as fake or real. The contingency table is represented using a color scale, with annotations showing the counts for each combination of popularity and label.

3. Pie Chart: Proportion of Fake/Real News for Popular vs Non-Popular News

Two pie charts are used to show the proportion of fake and real news articles within the popular and non-popular news categories. The pie charts make it easy to see the percentage of fake and real news in each category, helping to identify trends in news authenticity based on popularity.

Methodology for data preprocessing and feature extraction.

- **For Preprocessing the step's are:**
 - A function, `preprocess_text()`, is defined to clean and preprocess the text data:
 1. **Remove HTML Tags:** Eliminates any HTML markup from the text.
 2. **Remove Special Characters and Extra Whitespace:** Filters out non-alphanumeric characters and trims unnecessary spaces.
 3. **Normalize Text:** Converts all text to lowercase for uniformity.
 4. **Tokenize Text:** Splits the text into individual words using `word_tokenize()`.
 5. **Remove Stop Words:** Filters out common stop words (e.g., "and," "the") using the NLTK stopwords list.
- **Preprocessing Titles:**
 - The `preprocess_text` function is applied to the title column in the dataset.
 - A new column, `cleaned_title`, is created to store the preprocessed titles.
- **Save Processed Data:**
 - The cleaned dataset is saved as a CSV file, `combined_data_cleaned.csv`, for future use.
- **Verify Preprocessing:**
 - The first few rows of the title and `cleaned_title` columns are displayed to confirm the preprocessing results.

Before Preprocessing the `title` **looks like this :**

Example : -

Paris Jackson & Cara Delevingne Enjoy Night Out In Matching Outfits: They Have â€˜Amazing Chemistryâ€™™

After Preprocessing the `cleaned_title` **looks like this :**

paris jackson cara delevingne enjoy night matching outfits amazing chemistry

The `filter_csv_by_id` function is designed to clean and filter a CSV file based on multiple criteria:

1. **Remove Empty Cells:** Rows with any empty cell in any column are removed.

2. Remove Excessively Large Data: Rows where any cell contains text data exceeding a defined character length (1000 by default) are removed.
 3. Remove Invalid Tweet IDs: Rows where the 'tweet_ids' column is empty, null, or contains the value 0 are excluded.
 4. Filter Rows by ID: Only rows where the 'id' column starts with a letter are kept.
 5. Rename Column: The 'id' column is renamed to 'news_id' for better clarity.
- Execution Steps:
 - The script reads the input CSV file from the specified input_path.
 - It applies the filtering conditions and creates a new filtered dataset.
 - The filtered data is saved to a new CSV file (output_path).
 - Statistics:
 - The script prints statistics on how many rows were kept and removed during the filtering process.
 - Error Handling:
 - If the input file is not found, an error message is displayed.
 - General errors are caught and their details are printed.
-
- Original number of rows: 23196
 - Rows kept: 13110
 - Rows removed: 10086
 - New file saved as: dataset/filtered_data.csv

Feature Extraction in the Project

Feature extraction is a critical step in this project, as it transforms raw data into meaningful representations that the machine learning model can interpret. The features are derived from textual content, named entities, and engagement metrics to predict article popularity. Below is a detailed explanation of the feature extraction process used in this project:

1. Sentiment Score (Text Polarity)

- Description: The sentiment_score is calculated using the TextBlob library, which analyzes the polarity of the article title.
 - A positive score indicates positive sentiment.
 - A negative score indicates negative sentiment.

- A score of zero indicates neutral sentiment.
- Purpose: This feature helps capture the emotional tone of the title and its potential impact on audience engagement.

code

```
data['sentiment_score'] = data['title'].apply(lambda x: TextBlob(str(x)).sentiment.polarity)
```

2. Title Length

- Description: This feature measures the number of characters in the article title.
- Purpose: Title length can influence user engagement; for example, shorter, concise titles might grab attention more effectively than longer ones.

code

```
data['title_length'] = data['title'].apply(len)
```

3. Word Count

- Description: This feature calculates the number of words in the article title.
- Purpose: Word count provides an additional measure of the article title's verbosity, which might affect how users perceive and interact with the content.

code

```
data['word_count'] = data['title'].apply(lambda x: len(str(x).split()))
```

4. Named Entities

- Description: The project extracts named entities such as person_name, org_name, and locations_name from the dataset. These are combined into a single textual feature, combined_entities.
- Purpose: Named entities represent key elements in the article and can influence its popularity. For example, articles mentioning well-known individuals or organizations might attract more attention.
- Implementation:
 - Missing values in the named entity columns are filled with placeholders (empty).
 - Named entities are concatenated into a unified textual feature.

code

```
data['combined_entities'] = (
    data['person_name'] + " " + data['org_name'] + " " + data['locations_name']
)
```

5. TF-IDF Vectorization

- **Description:** The TfidfVectorizer is used to transform textual data (title and combined_entities) into numerical features by calculating the Term Frequency-Inverse Document Frequency (TF-IDF) scores for a fixed set of 5,000 features.
- **Purpose:** TF-IDF helps capture the significance of words in the text while reducing the influence of commonly occurring words (e.g., "the", "and"). This is crucial for distinguishing between articles based on unique or impactful words.
- **Implementation:** The title and combined entities are vectorized, and the resulting sparse matrix is converted into a dense array for use in the model.

code

```
tfidf = TfidfVectorizer(max_features=5000)
X_title = tfidf.fit_transform(data['title'] + " " + data['combined_entities'])
```

6. Engagement Metrics

- **Description:** The dataset includes numerical features such as:
 - likes: Number of likes the article received.
 - shares: Number of times the article was shared.
 - comments: Number of comments on the article.
- **Purpose:** These features directly measure the popularity of the article and serve as key predictors in the model.

code

```
required_columns = ['likes', 'shares', 'comments']
```

7. Feature Integration

- After extracting individual features, they are combined into a single feature matrix:
 - TF-IDF Features: Representing textual information.
 - Numerical Features: Sentiment score, title length, word count, and engagement metrics.
- The features are concatenated to form the final input for the machine learning model.

code

```
X = np.hstack([
    X_title.toarray(),
    data[required_columns + ['sentiment_score', 'title_length', 'word_count']].values
])
```

Importance of Feature Extraction

- Improved Predictive Power: The features capture various aspects of the data, enabling the model to better understand relationships and patterns.
- Domain-Specific Insights: Features like named entities and sentiment scores provide domain-specific insights into article content and popularity.
- Reducing Noise: By transforming raw data into relevant features, the process filters out irrelevant information, improving model performance.

Details about the predictive modeling process and performance metrics.

Predictive Modeling Process and Performance Metrics

The predictive modeling process in this project is focused on classifying articles based on their popularity using a combination of textual and numerical features. The process encompasses data preparation, feature extraction, handling class imbalance, training a robust model, and evaluating its performance. Below are the detailed steps:

Predictive Modeling Process

1. Data Preparation

- **Dataset Loading:** The dataset is loaded from a CSV file, and missing values in key columns (person_name, org_name, and locations_name) are handled by replacing them with the placeholder 'empty'.
- **Feature Engineering:** Features such as sentiment score, title length, word count, and named entity combinations are created to enrich the dataset with meaningful information.
- **Text Vectorization:** The textual data (title and combined_entities) is transformed into numerical features using the TF-IDF vectorizer, creating a 5,000-dimensional feature space.

2. Handling Class Imbalance

- The dataset's target variable (label) is examined for class imbalance. Since imbalanced data can bias the model toward the majority class, Synthetic Minority Oversampling Technique (SMOTE) is applied to balance the classes in the training data.
- SMOTE generates synthetic samples of the minority class to ensure equal representation during training, improving the model's ability to predict both classes.

```
code
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

3. Model Selection

- A Random Forest Classifier is chosen for its ability to handle high-dimensional feature spaces and mixed types of data effectively. Additionally, Random Forest models are robust to overfitting due to their ensemble nature.
- Key hyperparameters include:
 - `n_estimators`: Number of trees in the forest (set to 200).
 - `max_depth`: Maximum depth of the trees (set to 30).
 - `class_weight`: Balanced to account for any residual class imbalance.

code

```
rf_model = RandomForestClassifier(  
    class_weight='balanced',  
    random_state=42,  
    n_estimators=200,  
    max_depth=30,  
    min_samples_split=10,  
    min_samples_leaf=1  
)
```

4. Model Training

- The model is trained on the resampled dataset created using SMOTE, ensuring balanced representation of classes during training.
- Features include both TF-IDF-transformed textual features and numerical features (likes, shares, comments, sentiment_score, etc.).

5. Model Prediction

- Predictions are generated for both the training and test datasets.
- Probabilistic predictions (`predict_proba`) are also obtained to evaluate performance using metrics such as the ROC-AUC score.

Performance Metrics

1. Accuracy

- Measures the proportion of correctly classified samples.
- Achieved Train Accuracy: 89.27% and Test Accuracy: 84.71%.

code

```
Random Forest - Train Accuracy: 0.8927
```

Random Forest - Test Accuracy: 0.8471

2. Precision, Recall, and F1-Score

- Precision: The proportion of true positives out of all predicted positives.
- Recall: The proportion of true positives out of all actual positives.
- F1-Score: The harmonic mean of precision and recall, balancing both metrics.
- Classification report for test data:

code

```
precision recall f1-score support

0 0.92 0.76 0.83 2271
1 0.80 0.93 0.86 2332

accuracy 0.85 4603
macro avg 0.86 0.85 0.85 4603
weighted avg 0.86 0.85 0.85 4603
```

3. Confusion Matrix

- Provides insights into the number of true positives, true negatives, false positives, and false negatives.
- Test Data Confusion Matrix:

code

```
[[1721 550] # Class 0: True Negatives and False Positives
 [ 154 2178] # Class 1: False Negatives and True Positives
```

4. ROC-AUC Score

- Evaluates the model's ability to distinguish between the two classes.
- A score close to 1 indicates excellent performance.
- Achieved ROC-AUC Score: 0.9149.

code

```
roc_auc = roc_auc_score(y_test, y_proba)
print(f"ROC-AUC Score: {roc_auc:.4f}")
```


Key Observations

1. High Precision and Recall:
 - The model demonstrates high precision and recall for both classes, ensuring balanced performance.
 - It is particularly effective at identifying Class 1 (popular articles) with a recall of 93%.
2. Overfitting Mitigation:
 - Although the model achieves a slightly higher accuracy on the training set (89.27%), the test accuracy (84.71%) indicates that overfitting is well-controlled.
3. Balanced Class Representation:
 - The use of SMOTE addresses class imbalance, enabling the model to perform well on minority class predictions.
4. Robustness of Random Forest:
 - The ensemble-based approach ensures robustness and reduces sensitivity to noise in the data.

Conclusion

The predictive modeling process effectively combines feature engineering, handling class imbalance, and leveraging the Random Forest algorithm to achieve a well-balanced, high-performing classifier. Performance metrics such as accuracy, precision, recall, F1-score, and ROC-AUC indicate the model's ability to classify articles based on their popularity with high reliability.

Insights on how named entities impact article engagement and popularity.

Insights on How Named Entities Impact Article Engagement and Popularity

Named entities, such as person names, organization names, and location names, play a crucial role in driving article engagement and popularity. Analyzing their relationships with metrics like likes, shares, and comments provides valuable insights into how audiences interact with content. Here's a breakdown of key insights from the data:

1. Named Entity Frequency and Engagement

- Articles featuring well-known person names (e.g., celebrities, public figures) tend to receive higher engagement. This is because audiences are more likely to interact with content related to familiar or influential individuals.
- Organization names, especially those of prominent brands, companies, or institutions, can attract specific audience segments. For example, articles mentioning tech giants like "Google" or "Apple" often spark interest in tech-savvy readers.
- Location names tied to trending events, travel, or crises (e.g., "Paris" for fashion, "California" for wildfires) generate engagement as readers seek information or updates.

Example Visualization: A bar chart showing the frequency of top-named entities and their corresponding average engagement metrics (likes, shares, comments).

2. Sentiment Polarity of Articles with Named Entities

- Named entities influence the sentiment score of articles, which impacts their popularity:
 - Positive sentiment: Articles about inspiring individuals or organizations performing good deeds tend to garner positive engagement (e.g., "Elon Musk announces breakthrough in renewable energy").
 - Negative sentiment: Controversial or tragic events associated with named entities also draw attention but may lead to polarized interactions.
- Articles with neutral sentiment and strong named entities, like news reports, can also achieve high engagement due to the inherent importance of the information.

Example Visualization: Scatter plot showing the correlation between named entity sentiment polarity and engagement metrics.

3. Relationship Between Combined Entities and Popularity

- Combining named entities (e.g., person + organization + location) can amplify an article's relevance. For instance:
 - "Tim Cook (person) announces new iPhone (organization) launch in San Francisco (location)" can generate higher engagement by appealing to multiple interests (tech, business, geography).
- Articles with rich entity combinations often outperform those with standalone entities, as they cater to broader audience interests.

Example Visualization: Heatmap showing the combined effect of entity counts on likes, shares, and comments.

4. Influence of Entity Types on Engagement Metrics

- Different entity types impact specific engagement metrics differently:
 - Likes: Strongly influenced by well-known individuals or organizations.
 - Shares: More likely when entities are tied to actionable content, such as causes or trending topics (e.g., "Red Cross efforts in Ukraine").
 - Comments: Spurred by controversial entities or those tied to polarizing topics, leading to discussions or debates.

Example Visualization: A stacked bar chart showing the distribution of engagement metrics across entity types.

5. Entity Popularity Trends Over Time

- Articles mentioning certain named entities gain more popularity during specific periods. For example:
 - Seasonal topics: "New York" during New Year celebrations or "Amazon" during holiday sales.
 - Event-driven spikes: Named entities like "India" and "World Cup" experience peaks during major sports tournaments.
- These trends highlight the temporal relevance of entities in driving engagement.

Example Visualization: A time-series plot showing fluctuations in engagement for top-named entities.

6. Named Entities and Virality

- Articles with unique or rarely mentioned named entities (e.g., niche organizations or lesser-known locations) often attract niche audiences, leading to organic sharing within specific communities.
- On the other hand, articles featuring globally recognized entities (e.g., "Taylor Swift," "NASA") tend to go viral due to widespread interest.

Example Insight: Viral potential is higher for articles combining trending named entities with strong emotional or informational appeal.

Actionable Insights

- **Content Strategy:** Focus on incorporating trending or high-interest named entities in article titles and content to boost engagement.
- **Target Audience:** Use specific entities to cater to niche audiences, ensuring the content resonates with their interests.
- **Sentiment Optimization:** Balance the sentiment associated with named entities to maximize audience reactions, whether it's inspiring positivity or fostering discussions on critical issues.

By analyzing the role of named entities in article engagement, content creators can optimize their strategies to capture audience interest, increase visibility, and drive meaningful interactions.